

Sales Prediction: The Dataset used in these models tells about whether a person of certain age having certain income purchases a product or not. We need to predict whether a targeted audience will purchase the product or not.

```
In [1]: # import the library
import pandas as pd
# load the dataset
dataset = pd.read_csv('Social_Network_Ads.csv')

In [2]: # display the first five rows
dataset.head()

Out[2]:   User ID  Gender  Age  EstimatedSalary  Purchased
0    15624510    Male    19        19000         0
1    15810944    Male    35        20000         0
2    15668575  Female    26        43000         0
3    15603246  Female    27        57000         0
4    15804002    Male    19        76000         0

In [3]: # check the null values
dataset.isnull().sum()

Out[3]: User ID      0
Gender       0
Age          0
EstimatedSalary  0
Purchased     0
dtype: int64

In [4]: # ID is not helpful in predicting , so we will drop ID
dataset.drop(columns=['User ID'],inplace = True)

In [5]: # Define X and y
# X is independent variable/features (Gender , Age and Estimated Salary)
# y is dependent variable/label/target (Purchased)
X = dataset.iloc[:, 0:3]
y = dataset.iloc[:, 3]

In [6]: # Character values in gender to be converted to numeric
# import the library
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder() # instantiating object of Labelencoder class.
X['Gender'] = le.fit_transform(X['Gender'])
X.head()

Out[6]:   Gender  Age  EstimatedSalary
0      1    19        19000
1      1    35        20000
2      0    26        43000
3      0    27        57000
4      1    19        76000

In [7]: # Split into train and test
# train will be used for training and test for testing
from sklearn.model_selection import train_test_split
X_train,X_test, y_train, y_test = train_test_split(X, y , test_size=0.2 , random_state=5)

In [8]: dataset.describe()

Out[8]:   Age  EstimatedSalary  Purchased
count  400.000000  400.000000  400.000000
mean   37.655000  69742.500000  0.357500
std    10.482877  34096.960282  0.479864
min    18.000000  15000.000000  0.000000
25%   29.750000  43000.000000  0.000000
50%   37.000000  70000.000000  0.000000
75%   46.000000  88000.000000  1.000000
max    60.000000  150000.000000  1.000000

In [9]: dataset.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 4 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Gender            400 non-null    object  
 1   Age               400 non-null    int64  
 2   EstimatedSalary   400 non-null    int64  
 3   Purchased         400 non-null    int64  
dtypes: int64(3), object(1)
memory usage: 12.6+ KB

In [10]: # We need to standardise the values
# import the library
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

In [12]: # Build the model
# import the library
import tensorflow as tf
from keras.layers import *
from keras.models import *

In [13]: model = Sequential()
# input layer
model.add(Dense(16,input_dim=3,activation='relu'))
# hidden layer
model.add(Dense(16,activation='relu'))
model.add(Flatten())
# output layer
model.add(Dense(1 , activation= 'sigmoid'))
model.summary()

Model: "sequential"


```

Layer (type) Output Shape Param #
=====
dense (Dense) (None, 16) 64
dense_1 (Dense) (None, 16) 272
flatten (Flatten) (None, 16) 0
dense_2 (Dense) (None, 1) 17
=====
Total params: 353 (1.38 KB)
Trainable params: 353 (1.38 KB)
Non-trainable params: 0 (0.00 Byte)

```



In [14]: # Compile the model
model.compile(loss='binary_crossentropy',optimizer='adam',metrics='accuracy')

In [15]: # use fit to train the model
model.fit(X_train,y_train,batch_size=16,epochs=10)

Epoch 1/10
20/20 [=====] - 1s 3ms/step - loss: 0.7318 - accuracy: 0.2344
Epoch 2/10
20/20 [=====] - 0s 3ms/step - loss: 0.6931 - accuracy: 0.4750
Epoch 3/10
20/20 [=====] - 0s 2ms/step - loss: 0.6577 - accuracy: 0.7375
Epoch 4/10
20/20 [=====] - 0s 2ms/step - loss: 0.6208 - accuracy: 0.8156
Epoch 5/10
20/20 [=====] - 0s 2ms/step - loss: 0.5818 - accuracy: 0.8625
Epoch 6/10
20/20 [=====] - 0s 2ms/step - loss: 0.5429 - accuracy: 0.8750
Epoch 7/10
20/20 [=====] - 0s 2ms/step - loss: 0.5011 - accuracy: 0.8719
Epoch 8/10
20/20 [=====] - 0s 2ms/step - loss: 0.4613 - accuracy: 0.8781
Epoch 9/10
20/20 [=====] - 0s 2ms/step - loss: 0.4226 - accuracy: 0.8781
Epoch 10/10
20/20 [=====] - 0s 2ms/step - loss: 0.3872 - accuracy: 0.8844
Out[15]: <keras.src.callbacks.History at 0x78e7ddd75570>

In [19]: # predict the values
y_pred = model.predict(X_test)
y_pred = (y_pred>0.5)

3/3 [=====] - 0s 4ms/step

In [20]: from sklearn.metrics import confusion_matrix,accuracy_score
confusion_matrix(y_test,y_pred)

Out[20]: array([[49,  4],
   [ 2, 25]])

In [21]: accuracy_score(y_test,y_pred)

Out[21]: 0.925
```