

▼ RNN using MNIST

- MNIST data set can be treated as sequence of rows and columns of pixels.
- Process MNIST image as 28-element input vector and timesteps equal to 28.

```
# import the library
import numpy as np
from keras.models import Sequential
from keras.layers import Dense, Activation , SimpleRNN
from keras.utils import to_categorical
from keras.datasets import mnist

# load mnist dataset
(x_train,y_train), (x_test,y_test) = mnist.load_data()
print(x_train.shape)

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 [=====] - 0s 0us/step
(60000, 28, 28)

# compute the number of labels
num_labels = len(np.unique(y_train))
num_labels

10

# convert to categorical
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)

# normalise
x_train = x_train.astype('float')/255
x_test = x_test.astype('float')/255

# network parameters
image_size= x_train.shape[1]
input_shape= (28,28)
batch_size= 128
units = 32

# create simplernn model
model = Sequential()
model.add(SimpleRNN(units=units,input_shape=input_shape))
model.add(Dense(num_labels))
model.add(Activation('softmax'))
model.summary()
```

```
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
=====		
simple_rnn_1 (SimpleRNN)	(None, 32)	1952
dense_1 (Dense)	(None, 10)	330
activation_1 (Activation)	(None, 10)	0
=====		
Total params: 2282 (8.91 KB)		
Trainable params: 2282 (8.91 KB)		
Non-trainable params: 0 (0.00 Byte)		

```
# compile the model , define loss function , optimiser , metrics
model.compile(loss='categorical_crossentropy', optimizer='sgd',metrics=['accuracy'])

# train the model
model.fit(x_train,y_train,epochs=10,batch_size=batch_size)

Epoch 1/10
469/469 [=====] - 5s 12ms/step - loss: 0.6316 - accuracy: 0.
Epoch 2/10
469/469 [=====] - 3s 7ms/step - loss: 0.5797 - accuracy: 0.8
Epoch 3/10
469/469 [=====] - 3s 7ms/step - loss: 0.5312 - accuracy: 0.8
Epoch 4/10
469/469 [=====] - 4s 8ms/step - loss: 0.5028 - accuracy: 0.8
Epoch 5/10
469/469 [=====] - 5s 10ms/step - loss: 0.4820 - accuracy: 0.
Epoch 6/10
469/469 [=====] - 3s 7ms/step - loss: 0.4552 - accuracy: 0.8
Epoch 7/10
469/469 [=====] - 3s 7ms/step - loss: 0.4457 - accuracy: 0.8
Epoch 8/10
469/469 [=====] - 4s 9ms/step - loss: 0.4285 - accuracy: 0.8
Epoch 9/10
469/469 [=====] - 4s 9ms/step - loss: 0.4127 - accuracy: 0.8
Epoch 10/10
469/469 [=====] - 3s 7ms/step - loss: 0.3961 - accuracy: 0.8
<keras.src.callbacks.History at 0x7fea1f73a740>
```



```
# Evaluate the model
```

```
loss , acc = model.evaluate(x_test,y_test,batch_size=batch_size)
```

```
79/79 [=====] - 0s 4ms/step - loss: 0.3772 - accuracy: 0.894
```



```
print('Test Accuracy :', round(100*acc,2))
```

```
Test Accuracy : 89.43
```

```
#predict
pred = model.predict(x_test[:2])

1/1 [=====] - 0s 21ms/step

# predicted value
pred[1].argmax()

2

# actual value
y_test[1].argmax()

2

y_test[1]

array([0., 0., 1., 0., 0., 0., 0., 0., 0.], dtype=float32)

y_test[0]

➡ array([0., 0., 0., 0., 0., 0., 0., 1., 0., 0.], dtype=float32)
```